

Loan Approval Prediction

May 23, 2023

1 Importing Necessary Libraries

```
[1]: import pandas as pd, numpy as np # For Data Manipulation
import matplotlib.pyplot as pp, seaborn as sb # For Data Visualisation
from category_encoders import OrdinalEncoder # For Categorical Encoding
from sklearn.model_selection import train_test_split # For Model Validation sets
from sklearn.naive_bayes import GaussianNB # Model 1 : Naive Bayes ->
    Probabilistic Model
from sklearn.tree import DecisionTreeClassifier as DTC # Model 2 : Decision
    Tree - Deterministic Model
import sklearn.tree as tree # For Decision Tree Visualisation
from sklearn.ensemble import RandomForestClassifier # Model 3 : Random Forest ->
    Ensemble Model (Bagging)
from sklearn.metrics import
    accuracy_score, classification_report, roc_curve, roc_auc_score # Model
    Validation
```

2 Importing our data

```
[2]: data = pd.read_csv('data.csv')
```

3 Exploring our data - Descriptive

3.1 Getting first 5 rows

```
[3]: data.head()
```

```
[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	

1	4583	1508.0	128.0	360.0
2	3000	0.0	66.0	360.0
3	2583	2358.0	120.0	360.0
4	6000	0.0	141.0	360.0

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

3.2 Getting first 'n' rows

```
[4]: data.head(10)
```

```
[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
5	LP001011	Male	Yes	2	Graduate	Yes	
6	LP001013	Male	Yes	0	Not Graduate	No	
7	LP001014	Male	Yes	3+	Graduate	No	
8	LP001018	Male	Yes	2	Graduate	No	
9	LP001020	Male	Yes	1	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
5	5417	4196.0	267.0	360.0	
6	2333	1516.0	95.0	360.0	
7	3036	2504.0	158.0	360.0	
8	4006	1526.0	168.0	360.0	
9	12841	10968.0	349.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
5	1.0	Urban	Y

6	1.0	Urban	Y
7	0.0	Semiurban	N
8	1.0	Urban	Y
9	1.0	Semiurban	N

3.3 Getting last 5 rows

```
[5]: data.tail()
```

```
[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	Credit_History	Property_Area	Loan_Status
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

3.4 Getting last 'n' rows

```
[6]: data.tail(10)
```

```
[6]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
604	LP002959	Female	Yes	1	Graduate	No	
605	LP002960	Male	Yes	0	Not Graduate	No	
606	LP002961	Male	Yes	1	Graduate	No	
607	LP002964	Male	Yes	2	Not Graduate	No	
608	LP002974	Male	Yes	0	Graduate	No	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
--	-----------------	-------------------	------------	------------------	---

604	12000	0.0	496.0	360.0
605	2400	3800.0	NaN	180.0
606	3400	2500.0	173.0	360.0
607	3987	1411.0	157.0	360.0
608	3232	1950.0	108.0	360.0
609	2900	0.0	71.0	360.0
610	4106	0.0	40.0	180.0
611	8072	240.0	253.0	360.0
612	7583	0.0	187.0	360.0
613	4583	0.0	133.0	360.0

	Credit_History	Property_Area	Loan_Status
604	1.0	Semiurban	Y
605	1.0	Urban	N
606	1.0	Semiurban	Y
607	1.0	Rural	Y
608	1.0	Rural	Y
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

3.5 Looking at the description

```
[7]: data.describe()
```

```
[7]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.00000
mean	5403.459283	1621.245798	146.412162	342.00000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.00000
25%	2877.500000	0.000000	100.000000	360.00000
50%	3812.500000	1188.500000	128.000000	360.00000
75%	5795.000000	2297.250000	168.000000	360.00000
max	81000.000000	41667.000000	700.000000	480.00000

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

3.6 Looking at its shape

```
[8]: data.shape
```

```
[8]: (614, 13)
```

4 Data Preperation

4.1 Checking for Null Values

```
[9]: data.isnull().sum()
```

```
[9]: Loan_ID          0
     Gender          13
     Married         3
     Dependents      15
     Education        0
     Self_Employed   32
     ApplicantIncome  0
     CoapplicantIncome 0
     LoanAmount      22
     Loan_Amount_Term 14
     Credit_History   50
     Property_Area    0
     Loan_Status      0
     dtype: int64
```

4.2 Handling Null values

4.2.1 Handling Null Values in Categorical Features

```
[10]: data['Gender'].fillna(method='ffill',inplace=True)
      data['Self_Employed'].fillna(method='bfill',inplace=True)
      data['Married'].fillna(method='ffill',inplace=True)
```

```
[11]: data.isnull().sum()
```

```
[11]: Loan_ID          0
     Gender           0
     Married          0
     Dependents       15
     Education         0
     Self_Employed    0
     ApplicantIncome  0
     CoapplicantIncome 0
     LoanAmount       22
     Loan_Amount_Term 14
     Credit_History   50
```

```
Property_Area      0
Loan_Status        0
dtype: int64
```

4.2.2 Handling Null Values in Dependents by mean

```
[12]: data['Dependents'].value_counts()
```

```
[12]: 0      345
      1      102
      2      101
      3+      51
      Name: Dependents, dtype: int64
```

Converting the data into numerical

```
[13]: def convert_to_str(x):
      if x == '3+':
          return float(x[:-1])
      return float(x)
```

```
[14]: data['Dependents'] = data['Dependents'].apply(convert_to_str)
```

```
[15]: data['Dependents'].fillna(data['Dependents'].mean(), inplace=True)
```

```
[16]: data['Dependents'] = data['Dependents'].apply(float)
      data['Dependents'].isnull().sum()
```

```
[16]: 0
```

4.2.3 Handling Null Values in Numerical Features

```
[17]: data['LoanAmount'].fillna(data['LoanAmount'].mean(), inplace=True)
      data['Loan_Amount_Term'].fillna(data['Loan_Amount_Term'].mean(), inplace=True)
      data['Credit_History'].fillna(data['Credit_History'].mean(), inplace=True)
```

```
[18]: data.isnull().sum()
```

```
[18]: Loan_ID      0
      Gender      0
      Married     0
      Dependents  0
      Education   0
      Self_Employed  0
      ApplicantIncome  0
      CoapplicantIncome  0
      LoanAmount   0
      Loan_Amount_Term  0
```

```
Credit_History      0
Property_Area       0
Loan_Status         0
dtype: int64
```

4.3 Looking for duplicates

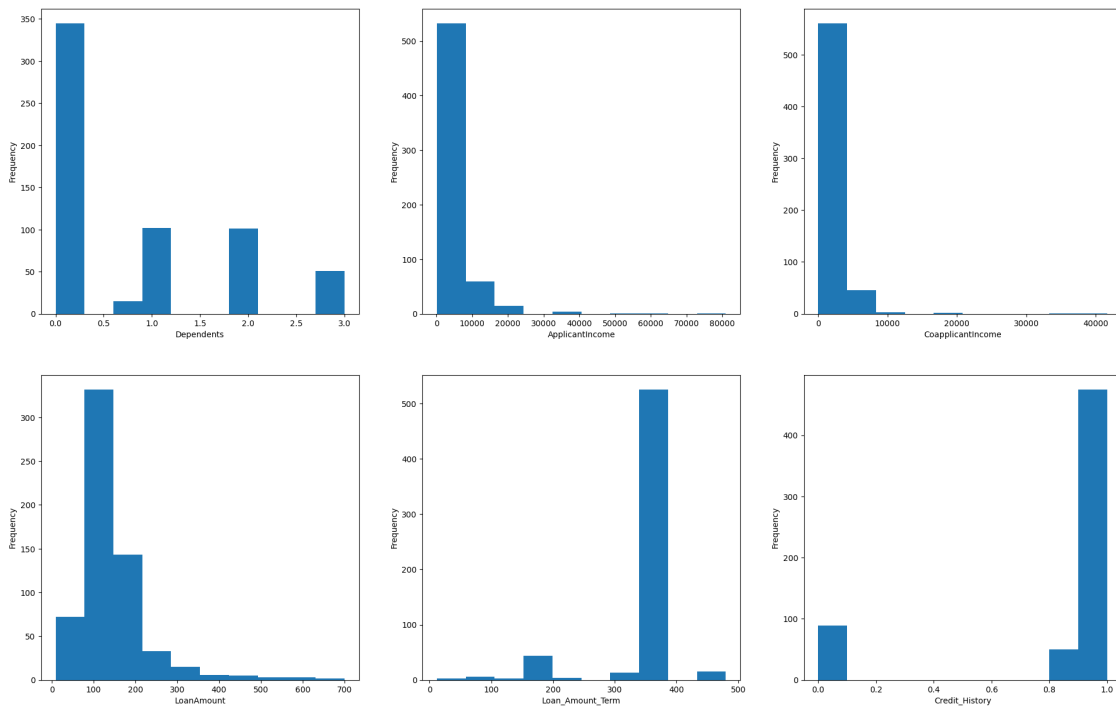
```
[19]: data.duplicated().sum()
```

```
[19]: 0
```

5 Exploring our data - Graphical

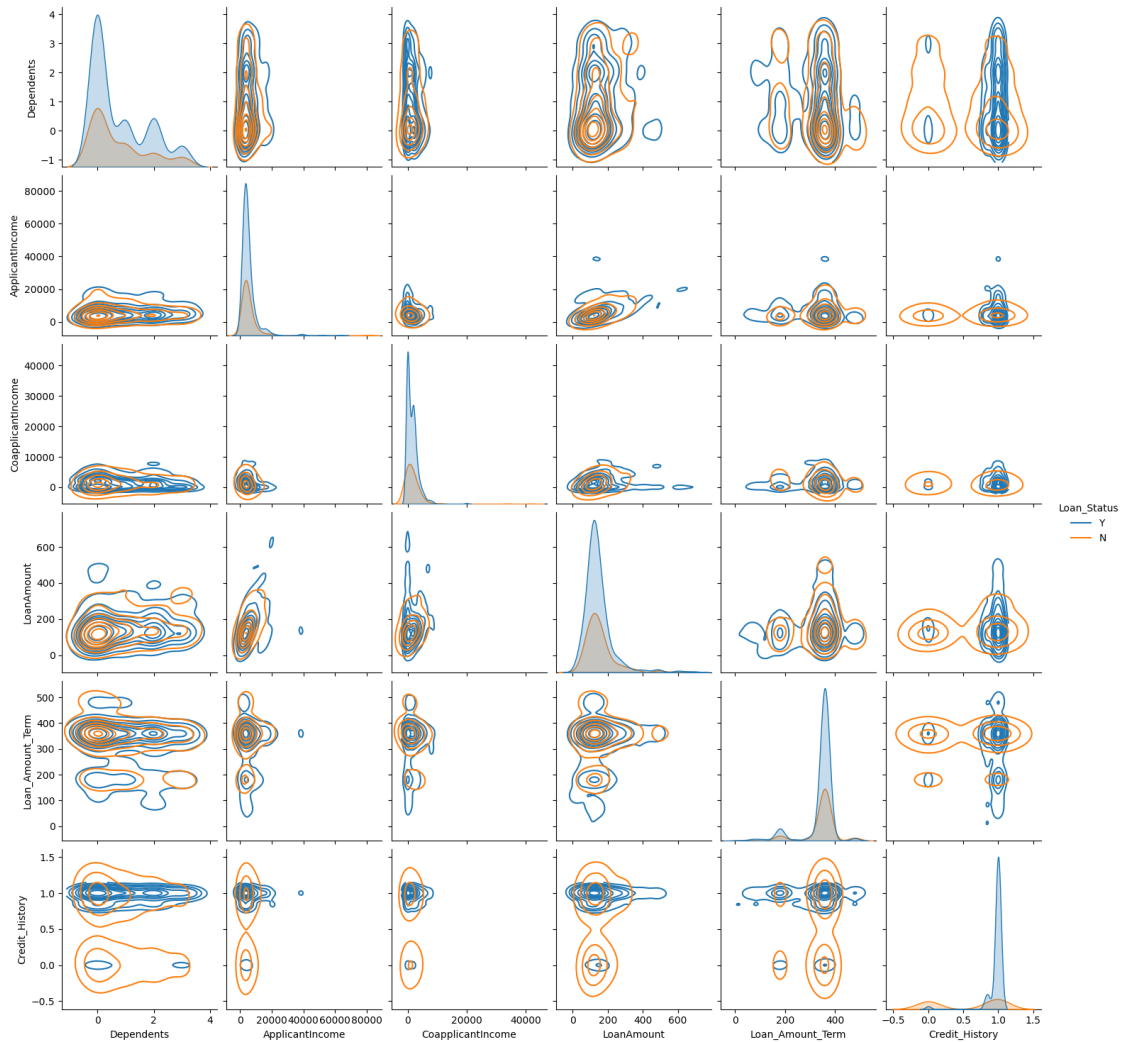
5.1 Histograms

```
[20]: numeric_vars = ['Dependents', 'ApplicantIncome', 'CoapplicantIncome',  
                    ↪ 'LoanAmount',  
                    'Loan_Amount_Term', 'Credit_History']  
# create histograms for each numeric variable  
fig = pp.figure(figsize=(24, 15))  
for i in range(len(numeric_vars)):  
    var = numeric_vars[i]  
    sub = fig.add_subplot(2, 3, i + 1)  
    sub.set_xlabel(var)  
    data[var].plot(kind = 'hist')
```



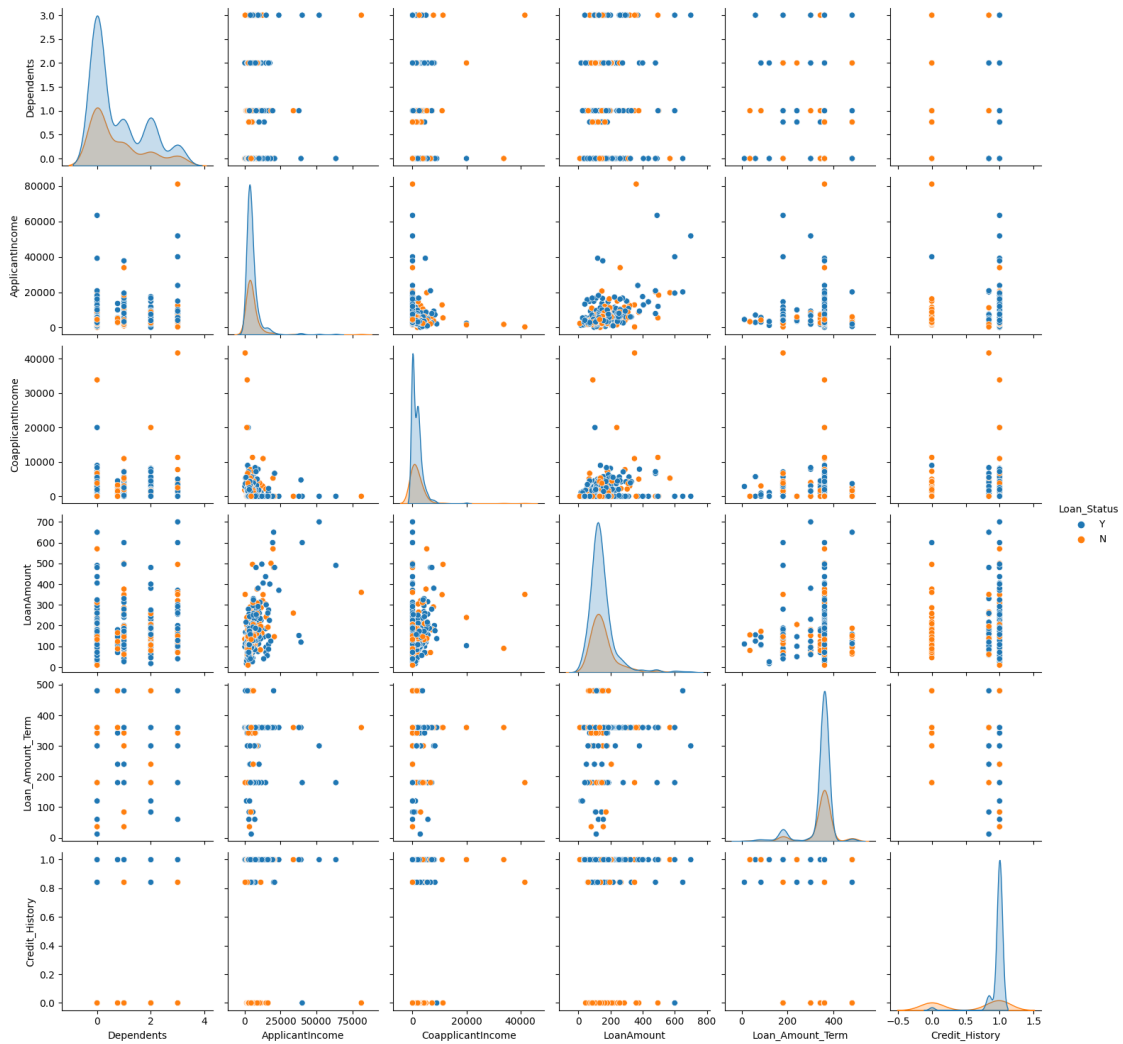
5.2 Kde Plot

```
[21]: sb.pairplot(data,hue='Loan_Status',kind='kde')  
pp.show()
```



5.3 Pair Plot

```
[22]: sb.pairplot(data,hue='Loan_Status')  
pp.show()
```

6 Encoding Categorical Features

[23]: `data.head()`

```
[23]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0.0	Graduate	No	
1	LP001003	Male	Yes	1.0	Graduate	No	
2	LP001005	Male	Yes	0.0	Graduate	Yes	
3	LP001006	Male	Yes	0.0	Not Graduate	No	
4	LP001008	Male	No	0.0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	146.412162	360.0	
1	4583	1508.0	128.000000	360.0	

2	3000	0.0	66.000000	360.0
3	2583	2358.0	120.000000	360.0
4	6000	0.0	141.000000	360.0

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

6.1 Encoding for non continuous features

```
[24]: data.drop('Loan_ID',axis=1,inplace=True)
```

6.2 Encoding other features

6.2.1 Looking at our categorical Features in detail

```
[25]: cf = {}
      ↪ ['Gender', 'Married', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']
      for i in cf:
          print('-----',i,'-----')
          print(data[i].value_counts())
          print('\n')
```

```
----- Gender -----
Male      500
Female    114
Name: Gender, dtype: int64
```

```
----- Married -----
Yes       400
No        214
Name: Married, dtype: int64
```

```
----- Education -----
Graduate      480
Not Graduate   134
Name: Education, dtype: int64
```

```
----- Self_Employed -----
No          528
Yes          86
Name: Self_Employed, dtype: int64
```

```

----- Property_Area -----
Semiurban    233
Urban        202
Rural        179
Name: Property_Area, dtype: int64

```

```

----- Loan_Status -----
Y    422
N    192
Name: Loan_Status, dtype: int64

```

6.2.2 Defining custom mapping

```

[26]: maps = [{'col': 'Gender',
               'mapping': {'Male': 1, 'Female': 0}},
              {'col': 'Married',
               'mapping': {'Yes': 1, 'No': 0}},
              {'col': 'Education',
               'mapping': {'Graduate': 1, 'Not Graduate': 0}},
              {'col': 'Self_Employed',
               'mapping': {'Yes': 1, 'No': 0}},
              {'col': 'Property_Area',
               'mapping': {'Rural': 0, 'Semiurban': 1, 'Urban': 2}},
              {'col': 'Loan_Status',
               'mapping': {'Y': 1, 'N': 0}}]

```

6.2.3 Encoding our data

```

[27]: oc = OrdinalEncoder(cols=cf, return_df=True,
                           mapping=maps)
data[cf] = oc.fit_transform(data[cf])

```

6.2.4 Verifying our data

```

[28]: data.head()

```

```

[28]:   Gender  Married  Dependents  Education  Self_Employed  ApplicantIncome  \
0      1         0         0.0         1         0           5849
1      1         1         1.0         1         0           4583
2      1         1         0.0         1         1           3000
3      1         1         0.0         0         0           2583
4      1         0         0.0         1         0           6000

```

	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	\
0	0.0	146.412162	360.0	1.0	
1	1508.0	128.000000	360.0	1.0	
2	0.0	66.000000	360.0	1.0	
3	2358.0	120.000000	360.0	1.0	
4	0.0	141.000000	360.0	1.0	

	Property_Area	Loan_Status
0	2	1
1	0	0
2	2	1
3	2	1
4	2	1

```
[29]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                 614 non-null   int32
1   Married                614 non-null   int32
2   Dependents             614 non-null   float64
3   Education              614 non-null   int32
4   Self_Employed          614 non-null   int32
5   ApplicantIncome        614 non-null   int64
6   CoapplicantIncome      614 non-null   float64
7   LoanAmount             614 non-null   float64
8   Loan_Amount_Term       614 non-null   float64
9   Credit_History         614 non-null   float64
10  Property_Area          614 non-null   int32
11  Loan_Status            614 non-null   int32
dtypes: float64(5), int32(6), int64(1)
memory usage: 43.3 KB
```

7 Splitting our data set

```
[30]: x = data.drop('Loan_Status',axis=1)
      y = data['Loan_Status']
      x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.33,
      random_state=42)
```

```
[31]: x_train.head()
```

```
[31]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	\
244	1	1	0.0	0	0	3406	
393	1	1	2.0	0	0	1993	
310	0	0	0.0	1	0	2917	
408	1	1	1.0	1	0	8300	
572	1	1	2.0	1	0	16666	

	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	\
244	4417.0	123.0	360.0	1.0	
393	1625.0	113.0	180.0	1.0	
310	0.0	84.0	360.0	1.0	
408	0.0	152.0	300.0	0.0	
572	0.0	275.0	360.0	1.0	

	Property_Area
244	1
393	1
310	1
408	1
572	2

```
[32]: x_test.head()
```

```
[32]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	\
350	1	1	0.0	1	0	9083	
377	1	1	0.0	1	0	4310	
163	1	1	2.0	1	0	4167	
609	0	0	0.0	1	0	2900	
132	1	0	0.0	1	0	2718	

	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	\
350	0.0	228.0	360.0	1.000000	
377	0.0	130.0	360.0	0.842199	
163	1447.0	158.0	360.0	1.000000	
609	0.0	71.0	360.0	1.000000	
132	0.0	70.0	360.0	1.000000	

	Property_Area
350	1
377	1
163	0
609	0
132	1

```
[33]: y_train.head()
```

```
[33]: 244    1
      393    1
      310    1
      408    0
      572    1
      Name: Loan_Status, dtype: int32
```

```
[34]: y_test.head()
```

```
[34]: 350    1
      377    1
      163    1
      609    1
      132    1
      Name: Loan_Status, dtype: int32
```

8 Training our models

8.1 Model 1 : Naive Bayes (Gaussian Naive Bayes)

```
[35]: nb = GaussianNB()
      nb.fit(x_train,y_train)
```

```
[35]: GaussianNB()
```

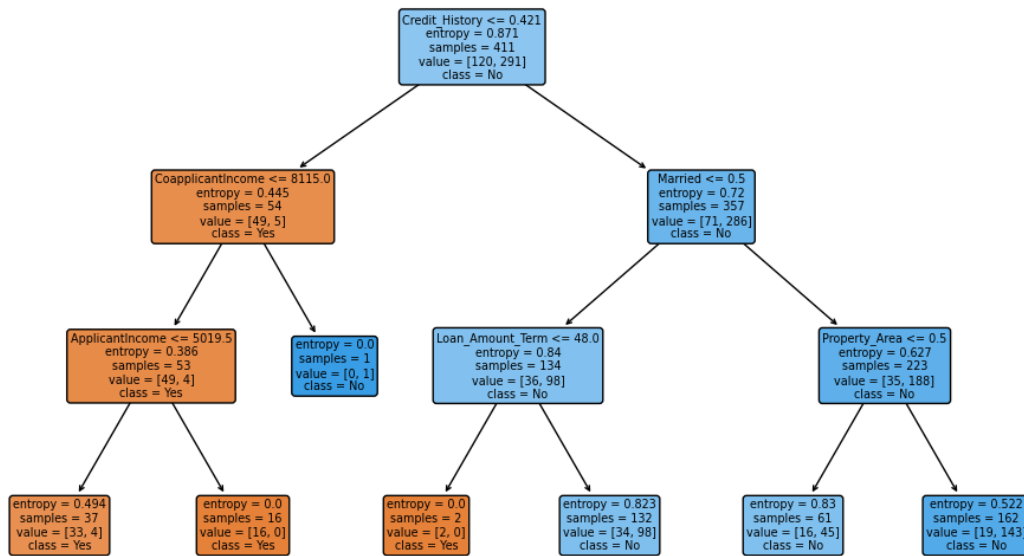
8.2 Model 2 : Decision Tree Classifier

```
[36]: dtc = DTC(criterion="entropy", max_depth=3)
      dtc.fit(x_train,y_train)
```

```
[36]: DecisionTreeClassifier(criterion='entropy', max_depth=3)
```

8.2.1 Visualisng the decision Tree

```
[37]: pp.figure(figsize=(12,7))
      tree.plot_tree(dtc.fit(x_train,y_train),feature_names = x_train.
      ↪columns,class_names=['Yes','No'],
      filled=True,rounded=True)
      pp.show()
```



8.3 Model 3 : Random Forest Classifier

```
[38]: rfc = RandomForestClassifier(criterion='entropy',max_depth=3,random_state=42,n_jobs=-1)
rfc.fit(x_train,y_train)
```

```
[38]: RandomForestClassifier(criterion='entropy', max_depth=3, n_jobs=-1,
                             random_state=42)
```

9 Evaluationg our models

9.1 Accuracy Scores

```
[39]: models = ["Naive Bayes","Decision Tree","Random Forest"]
acc = pd.DataFrame({
    "train" : [accuracy_score(y_train,nb.
        ↪predict(x_train))*100,accuracy_score(y_train,dtc.predict(x_train))*100,
              accuracy_score(y_train,rfc.predict(x_train))*100],
    "test" : [accuracy_score(y_test,nb.
        ↪predict(x_test))*100,accuracy_score(y_test,dtc.predict(x_test))*100,
              accuracy_score(y_test,rfc.predict(x_test))*100]
},index=models)
```

```
[40]: acc
```

```
[40]:
```

	train	test
Naive Bayes	81.265207	79.802956
Decision Tree	82.238443	79.310345
Random Forest	81.508516	79.802956

9.2 Classification Reports

9.2.1 Naive Bayes

```
[41]: print(classification_report(y_train,nb.  
    ↪predict(x_train),target_names=["Yes","No"]))
```

	precision	recall	f1-score	support
Yes	0.81	0.47	0.59	120
No	0.81	0.96	0.88	291
accuracy			0.81	411
macro avg	0.81	0.71	0.74	411
weighted avg	0.81	0.81	0.79	411

9.3 Decision Tree

```
[42]: print(classification_report(y_train,dtc.  
    ↪predict(x_train),target_names=["Yes","No"]))
```

	precision	recall	f1-score	support
Yes	0.93	0.42	0.58	120
No	0.81	0.99	0.89	291
accuracy			0.82	411
macro avg	0.87	0.71	0.74	411
weighted avg	0.84	0.82	0.80	411

9.4 Random Forest

```
[43]: print(classification_report(y_train,rfc.  
    ↪predict(x_train),target_names=["Yes","No"]))
```

	precision	recall	f1-score	support
Yes	0.91	0.41	0.56	120
No	0.80	0.98	0.88	291
accuracy			0.82	411
macro avg	0.85	0.70	0.72	411

weighted avg	0.83	0.82	0.79	411
--------------	------	------	------	-----

9.5 AUC ROC CURVE

9.6 Getting ROC Scores

```
[44]: pred_prob1 = nb.predict_proba(x_test)
      pred_prob2 = dtc.predict_proba(x_test)
      pred_prob3 = rfc.predict_proba(x_test)
      fpr1, tpr1, thresh1 = roc_curve(y_test, pred_prob1[:,1], pos_label=1)
      fpr2, tpr2, thresh2 = roc_curve(y_test, pred_prob2[:,1], pos_label=1)
      fpr3, tpr3, thresh3 = roc_curve(y_test, pred_prob3[:,1], pos_label=1)
      random_probs = [0 for i in range(len(y_test))]
      p_fpr, p_tpr, _ = roc_curve(y_test, random_probs, pos_label=1)
```

9.7 Getting AUC Scores

```
[45]: auc_score1 = roc_auc_score(y_test, pred_prob1[:,1])
      auc_score2 = roc_auc_score(y_test, pred_prob2[:,1])
      auc_score3 = roc_auc_score(y_test, pred_prob3[:,1])
      auc = pd.DataFrame(data=[auc_score1, auc_score2,
      ↪ auc_score3], index=models, columns=['AUC Score'])
      auc
```

```
[45]:
```

	AUC Score
Naive Bayes	0.736959
Decision Tree	0.735157
Random Forest	0.757740

9.8 Plotting the ROC Curve

```
[46]: pp.plot(fpr1, tpr1, linestyle='--', color='orange', label='Naive Bayes')
      pp.plot(fpr2, tpr2, linestyle='--', color='green', label='Decision Tree')
      pp.plot(fpr3, tpr3, linestyle='--', color='blue', label='Random Forest')
      pp.plot(p_fpr, p_tpr, linestyle='--', color='blue')
      pp.title('ROC curve')
      pp.xlabel('False Positive Rate')
      pp.ylabel('True Positive rate')
      pp.legend(loc='best')
      pp.savefig('ROC', dpi=300)
      pp.show()
```

